

PitchProX Functional Flow Cheat-Sheet (for backend & product teams — zero code)

Json.pages → Live-Call Engine

Step	Who/What	Action
Pre-Call	Front-end	Caller picks Lead Type (L001 ... L013) → POST to /session/start
During Call	STT stream	Every 3–5 s transcript snippet → /realtime/objection
Match Logic	Backend micro-service	1. Runs semantic search on objection_tree.statement
Prompt to Agent	WebSocket	Agent sees \leq 20-word reply in teleprompter
Trust_Cues	Backend only	Invisible counters that raise/lower intent score (0-1) every 15 s. Never exposed to caller. Used to decide:



Trust_Cues Explained

initial = starting intent (e.g., 25 for FSBO)

triggers = **if** elapsed seconds \geq t → add delta to intent

thresholds → map intent to **determine likelihood of successful call**

Impact: decides **when** to stop prompting at 7:00, **not** what to say.

Trust Score = rolling Intent Probability (0–1) updated every 15 s.

Data Sources

Sentiment – LLM classifies lead tone (-1 → +1).

Engagement Signals – word count, question ratio, filler words, interruptions.

Objection Layer Hit – deeper layers = lower baseline.

Positive Cues – micro-commitments (“sure”, “ok”, “tell me more”).

Negative Cues – silence > 3 s, abrupt topic shifts.

Algorithm Snapshot

Python

```
score = 0.5 # start neutral
score += sentiment * 0.3
score += engagement_signal * 0.4
score += positive_cues * 0.2
score -= negative_cues * 0.2
score = max(0, min(1, score))
```

Return the smoothed 3-second rolling average.

Sample Call Report → Post-Call Analytics

Step	Who/What	Action
Hang-Up	Backend	Collect raw transcript & meta
Scoring	LLM (Gemini) batch job	Runs sentiment, tone, trust deltas, prompt-adherence
Report	Auto-generated JSON	Stored at /calls/{call_id}/report
Display	Manager dashboard	Shows 87/100 score, emotional arc, pro-tips

After every call that has ended, a pop up screen pops up for the summary of the call, schedule of the follow up and the option to generate the call report. Yes, you create this report after every call, if the user requests it. However, each tier has a quota of how many will be included before prompt to pay for additional reports. The question I have is if the user does not generate it then, will they still be able to generate the report later through going back to the file. If so, then the option should be available in the call summary box at the end. I like the option for the call report at the end of the summary, because if there is ever a compliant we can review the call and see if it can aid in us helping go through it with the user.

Scoring is done by the LLM (Gemini) — not custom code.

Onboarding Calibration Engine → Agent Style Setter

Step	Who/What	Action
Day-Zero	Agent records 20 short test prompts	
Scoring	LLM grades: tone, pacing, word-choice, empathy	
Tier Chosen	Highest composite score → Empathy-High-Pacing, etc.	
Usage	Tier string is injected into every future prompt context so LLM speaks like the agent naturally talks	

Scores are calculated by the LLM
After calibration you simply pass tier to prompt context

Silent-7 Policy → Invisible Safety Net

Condition	Action
Wall-clock $\geq 7:00$ AND intent < 0.6	Backend mutes all AI prompts (agent keeps talking)
Wall-clock $\geq 7:00$ AND intent ≥ 0.6	Prompts continue (call is already winning)
User never hears “time’s up” or any credit prompt	

End-to-End Happy Path

Agent logs in → picks L001-Expired Listing

Calibration determined user was X from onboarding calibration → tier = Empathy-High-Pacing

Call starts → STT stream → objection matched → 20-word prompt delivered

Intent stays ≥ 0.6 → prompts flow until natural close

Hang-up → LLM auto-generates report → 87/100 score saved

Manager reviews → coaching tips surfaced

1. WHAT THE PLATFORM IS IN ONE SENTENCE

PitchProX is a **live, AI co-pilot for cold-calls** that whispers the *exact next sentence* into the agent's ear so they sound like a top-performer on every dial—without ever sounding canned.

2. THE THREE CORE MOMENTS OF TRUTH

Moment	Human Feeling	PitchProX Job
Pre-Call Nerves	"I hope I don't sound robotic"	Learn the agent's natural voice (Calibration)
Mid-Call Objection	"Uh-oh, what do I say now?"	Serve roughly a 20-word, emotionally-perfect reply (Live-Call Engine)
After Call	"Did I do well?"	Option for an after call report with coaching nuggets (Post-Call Analytics)

3. WALK-THROUGH OF EACH COMPONENT

A. ONBOARDING CALIBRATION

Why it exists:

Two agents can read the *same* sentence and one sounds like a friend, the other like a telemarketer.

Calibration fixes this:

How it works (step-by-step):

Prompt Series – The system plays 20 short scripts (e.g., *"this is the pricing point where I've done most of my business, which is why I would love to see the house"*)

Agent Reads Aloud during calibration when prompt when to say and how to say it – Voice is captured.

AI Grades – Speed, warmth, word-choice, pauses.

Tier Assigned – Empathy-High-Pacing, Clarity-Medium-Pacing, etc.

Invisible Forever – Every future prompt is *auto-rewritten* to match that style.

Developer note: store only the final tier string; raw audio is discarded after scoring.

B. LIVE-CALL ENGINE (Json.pages)

Why it exists:

Objections come in 0.8 seconds; humans need 2–3 seconds to think.

AI needs 0.2 seconds.

How it works:

Pick Lead Type – Agent taps "FSBO" or "Expired listing" before dialing.

Stream Transcript – STT sends real-time text every 3–5 s.

Match Objection – Vector search against `objection_tree`.

Prompt Teleprompter – Roughly 20-word reply appears; agent reads or ad-libs.

Trust Cues (Invisible) – Backend tallies intent score; if < 0.6 at 7:00 → prompts stop (Silent-7 minute rule policy). The indication is based on data on where or not the call will end in success, measured through the call to determine if the time nears 7 minutes should the AI continue to feed prompts or is this wasted prompting. We are looking to be effective and minimize our internal spend on calls that sit and talk talk talk for no actual benefit. This was if over 7 minutes and the call is not projected as a success, the call and user can continue to talk but with out prompts, so the user isn't upset and the call isn't expensive on out end for non success results. This prevents complaints.

C. POST-CALL ANALYTICS (Sample Call Report)

Why it exists:
User can learn why certain prompts were suggested and build on their skills. Overtime, they will be able to use what they are learning not just on the calls but during the in person appointments as well. Strengthening their overall skills. Also, managers can't listen to 100 calls/day, but they *can* skim a few-pages report scorecard.

How it works:
Call Ends – Raw transcript + meta data.
LLM Batch Job – Gemini scores sentiment, trust arc, prompt adherence.
Auto-Report – JSON drops into /calls/{id}/report.
Dashboard – 87/100, emotional arc, coaching tip.

4. THE SILENT-7 minute POLICY
Like an airbag: invisible unless needed.
Wall-clock hits 7:00 → if intent < 0.6 , *AI stops prompting*.
Wall-clock hits 7:00 → if intent > 0.6 , *AI continues prompting*.
No audible cue, no “time’s up.”
Call stays live; agent keeps talking naturally.

5. WHY THESE PIECES ARE NOT OPTIONAL

Piece	Remove It and...
Calibration	Prompts sound off-brand → CSAT drops 17 %.
Live-Call Engine	Agents revert to scripts → conversion rate halves.
Post-Call Report	Managers blind → coaching stops → churn rises.
Silent-7 minute policy	Pushes users to spend more time on the quality leads.

6. MENTAL ZIP FILE

Pre-Call → Calibrate voice
During → Whisper perfect lines
After → Get report coaching
Safety Net → Silent 7-minute cut-off (no one hears it)